

Spam Filter Performance Metrics Cheat Sheet

Accuracy:

$$\text{Accuracy} = (TP + TN) / (TP + TN + FP + FN)$$

Overall correctness: how many predictions were correct out of all predictions.

Precision (Positive Predictive Value):

$$\text{Precision} = TP / (TP + FP)$$

Out of all predicted spam emails, how many were actually spam.

Recall (Sensitivity, True Positive Rate):

$$\text{Recall} = TP / (TP + FN)$$

Out of all actual spam emails, how many were correctly identified as spam.

Specificity (True Negative Rate):

$$\text{Specificity} = TN / (TN + FP)$$

Out of all actual non-spam emails, how many were correctly identified as non-spam.

F1 Score:

$$F1 = 2 * (\text{Precision} * \text{Recall}) / (\text{Precision} + \text{Recall})$$

Harmonic mean of precision and recall balances both metrics.

False Positive Rate (Fall-out):

$$FPR = FP / (FP + TN)$$

Out of all actual non-spam emails, how many were incorrectly marked as spam.

1 - FPR (True Negative Rate / Specificity):

$$1 - FPR = TN / (TN + FP)$$

Also known as specificity; measures how well non-spam is recognized.

False Negative Rate (Miss Rate):

$$FNR = FN / (FN + TP)$$

Out of all actual spam emails, how many were missed (marked as non-spam).

ROC Curve:

Plot of TPR vs FPR at various threshold settings

Visual tool to assess trade-off between true positive and false positive rates.

AUC (Area Under Curve):

Area under the ROC curve (between 0 and 1)

Overall ability of the classifier to distinguish between classes.

NPV (Negative Predictive Value):

$$NPV = TN / (TN + FN)$$

Out of all predicted non-spam, how many were actually non-spam.

F2 Score:

$$F2 = (1 + 2^2) * (Precision * Recall) / (4 * Precision + Recall)$$

Gives more weight to recall useful when missing spam is costly.

F0.5 Score:

$$F0.5 = (1 + 0.5^2) * (Precision * Recall) / (0.25 * Precision + Recall)$$

Gives more weight to precision useful when false alarms are costly.

MCC (Matthews Correlation Coefficient):

$$MCC = (TP \cdot TN - FP \cdot FN) / \sqrt{(TP + FP)(TP + FN)(TN + FP)(TN + FN)}$$

Balanced score that works well even with imbalanced data.

Balanced Accuracy:

$$Balanced\ Accuracy = (Sensitivity + Specificity) / 2$$

Average of sensitivity and specificity useful in imbalanced data.

FDR (False Discovery Rate):

$$FDR = FP / (TP + FP)$$

Proportion of false spam predictions out of all predicted spam.

FOR (False Omission Rate):

$$FOR = FN / (TN + FN)$$

Proportion of missed spam out of all predicted non-spam.

Youdens J Statistic:

$$J = Sensitivity + Specificity - 1$$

Summarizes performance useful for finding optimal threshold.

Log Loss (Cross-Entropy Loss):

$$\text{LogLoss} = - (1/N) [y \log(p) + (1-y) \log(1-p)]$$

Penalizes incorrect probabilistic predictions more harshly.

Brier Score:

$$\text{Brier} = (1/N) (p - y)^2$$

Measures accuracy of probabilistic predictions lower is better.